

Yifan Zhong

Email : zhongyifan@stu.pku.edu.cn
Github: <https://github.com/Ivan-Zhong>
Mobile : (+86) 135-2130-5328

EDUCATION

- Peking University** Beijing, China
• Ph.D. candidate in Intelligence Science and Technology. Supervisor: Prof. Yaodong Yang. 2023 – 2028 (Expected)
- Tongji University** Shanghai, China
• B.S. in Computer Science and Technology. GPA: 4.96/5.00. Top 1%. 2019 – 2023

PUBLICATIONS

(* denotes equal contribution)

- Panacea: Pareto Alignment via Preference Adaptation for LLMs**
Yifan Zhong*, Chengdong Ma*, Xiaoyuan Zhang*, Ziran Yang, Haojun Chen, Qingfu Zhang, Siyuan Qi, Yaodong Yang
Under Review.
- Maximum Entropy Heterogeneous-Agent Reinforcement Learning**
Jiarong Liu*, Yifan Zhong*, Siyi Hu, Haobo Fu, QIANG FU, Xiaojun Chang, Yaodong Yang
International Conference on Learning Representations (ICLR), 2024. (Spotlight)
- CivRealm: A Learning and Reasoning Odyssey in Civilization for Decision-Making Agents**
Siyuan Qi*, Shuo Chen*, Yexin Li*, Xiangyu Kong*, Junqi Wang*, Bangcheng Yang, Pring Wong, Yifan Zhong,
Xiaoyuan Zhang, Zhaowei Zhang, Nian Liu, Yaodong Yang, Song-Chun Zhu
International Conference on Learning Representations (ICLR), 2024. (Spotlight)
- Heterogeneous-agent reinforcement learning**
Yifan Zhong*, Jakub Grudzien Kuba*, Xidong Feng*, Siyi Hu, Jiaming Ji, Yaodong Yang
Journal of Machine Learning Research (JMLR), 2024. (Github Star 360+)
- Safety gymnasium: A unified safe reinforcement learning benchmark**
Jiaming Ji*, Borong Zhang*, Jiayi Zhou*, Xuehai Pan, Weidong Huang, Ruiyang Sun, Yiran Geng, Yifan Zhong,
Juntao Dai, Yaodong Yang
Advances in Neural Information Processing Systems Track on Datasets and Benchmarks (NeurIPS DB), 2023.
- MARLLib: A Scalable and Efficient Multi-agent Reinforcement Learning Library**
Siyi Hu, Yifan Zhong, Minquan Gao, Weixun Wang, Hao Dong, Xiaodan Liang, Zhihui Li, Xiaojun Chang, Yaodong Yang
Journal of Machine Learning Research Machine Learning Open Source Software (JMLR MLOSS), 2023. (Github Star 780+)
- MyoChallenge 2022: Learning contact-rich manipulation using a musculoskeletal hand**
Vittorio Caggiano, Guillaume Durandau, Huwawei Wang, Alberto Chiappa, Alexander Mathis, Pablo Tano, Nisheet Patel, Alexandre Pouget, Pierre Schumacher, Georg Martius, Daniel Haeufle, Yiran Geng, Boshi An, Yifan Zhong,
Jiaming Ji, Yuanpei Chen, Hao Dong, Yaodong Yang, Rahul Siripurapu, Luis Eduardo Ferro Diez, Michael Kopp, Vihang Patil, Sepp Hochreiter, Yuval Tassa, Josh Merel, Randy Schultheis, Seungmoon Song, Massimo Sartori, Vikash Kumar
NeurIPS 2022 Competition Track.

WORKING EXPERIENCE

- Multi-Agent Systems Laboratory, BIGAI** Beijing, China
• Research Intern. Supervisor: Yaodong Yang, Siyuan Qi Aug. 2022 - Aug. 2024
Topic: Reinforcement Learning, Multi-Agent Systems, LLM Alignment

MISCELLANEOUS

- **Research Interests:** Reinforcement Learning, Large Language Model
- **Conference Reviewer:** NeurIPS (2024).

SELECTED HONORS & AWARDS

- **Second Prize of BIGAI Research Rising Star Award** 2024
- **Outstanding Graduate of Shanghai** 2023
- **First Place in NeurIPS 2022 Challenge Track** 2022
- **CCF Outstanding Undergraduate Student** 2022
- **National Scholarship** 2022, 2020
- **The 46th ICPC Asia Regional Contest (Jinan Site) Silver Medal** 2021
- **Qidi Scholarship** 2021
- **First Prize of Undergraduate Outstanding Student Scholarship** 2021